

# Aviation Accident Database & Synopses, up to 2023

## Introduction.

### Business understanding

A company is embarking on a strategic expansion into the aviation sector, aiming to diversify its investment portfolio by purchasing and operating planes for both commercial and private enterprises. As this venture is new to the organisation, there is limited internal knowledge about the aviation domain, particularly regarding safety and operational risk.

### Business problem

The company must identify which aircraft models and types have the lowest operational risk to guide safe and cost-effective entry into the aviation market.

### EDA objectives

This notebook explores aircraft risk patterns to support the Company's on safe aircraft acquisition decisions. The following objectives are structured around key EDA phases:

#### 1. Data understanding

- Load and preview the dataset to understand its structure, size and key variables.
- Identify the most relevant columns for analyzing aviation safety.

#### 2. Data cleaning

- Handle missing values and duplicated values.
- Check on inconsistent formatting in the key fields.
- Standardize categorical variables for accurate grouping.
- Check on outliers.....

#### 3. Univariate & Bivariate Analysis

- Explore the most common aircraft types and count how many incidents each has.
- Analyze incident trends over time to detect rising or declining risk patterns

#### 4. Multivariate & Grouped Insights

- Compare incident frequencies by usage type (commercial vs private), age, or manufacturer.

#### 5. Incident Causes & Locations

- Visualize the top causes of incidents and the regions with the highest incident density.

### 1. Data Understanding

- **Import the necessary libraries required for the data.**
- **Load the `Aviation.csv` data**

```
In [1]: # import the necessary libraries.  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [3]: # Load the data.  
df= pd.read_csv('AviationData.csv', encoding = 'latin1', low_memory = False)
```

```
In [5]: # Check the size of the data  
df.shape
```

```
Out[5]: (88889, 31)
```

```
In [7]: # check the columns of the data  
df.columns
```

```
Out[7]: Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date',  
              'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',  
              'Airport.Name', 'Injury.Severity', 'Aircraft.damage',  
              'Aircraft.Category', 'Registration.Number', 'Make', 'Model',  
              'Amateur.Built', 'Number.ofEngines', 'Engine.Type', 'FAR.Description',  
              'Schedule', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',  
              'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',  
              'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',  
              'Publication.Date'],  
             dtype='object')
```

```
In [9]: # check the whole information of the data  
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Event.Id                             88889 non-null  object
1   Investigation.Type                    88889 non-null  object
2   Accident.Number                      88889 non-null  object
3   Event.Date                           88889 non-null  object
4   Location                             88837 non-null  object
5   Country                             88663 non-null  object
6   Latitude                             34382 non-null  object
7   Longitude                            34373 non-null  object
8   Airport.Code                         50132 non-null  object
9   Airport.Name                         52704 non-null  object
10  Injury.Severity                      87889 non-null  object
11  Aircraft.damage                      85695 non-null  object
12  Aircraft.Category                    32287 non-null  object
13  Registration.Number                 87507 non-null  object
14  Make                                88826 non-null  object
15  Model                               88797 non-null  object
16  Amateur.Built                       88787 non-null  object
17  Number.of.Engines                   82805 non-null  float64
18  Engine.Type                         81793 non-null  object
19  FAR.Description                     32023 non-null  object
20  Schedule                            12582 non-null  object
21  Purpose.of.flight                   82697 non-null  object
22  Air.carrier                         16648 non-null  object
23  Total.Fatal.Injuries                 77488 non-null  float64
24  Total.Serious.Injuries               76379 non-null  float64
25  Total.Minor.Injuries                 76956 non-null  float64
26  Total.Uninjured                     82977 non-null  float64
27  Weather.Condition                   84397 non-null  object
28  Broad.phase.of.flight                61724 non-null  object
29  Report.Status                       82505 non-null  object
30  Publication.Date                     75118 non-null  object
dtypes: float64(5), object(26)
memory usage: 21.0+ MB

```

```

In [11]: # check the first five rows
df.head()

```

Out[11]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
0	20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States
1	20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States
3	20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States
4	20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States

5 rows × 31 columns

In [13]:

```
# check the last 5 rows
df.tail()
```

Out[13]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
88884	20221227106491	Accident	ERA23LA093	2022-12-26	Annapolis, MD	United States
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	United States
88886	20221227106497	Accident	WPR23LA075	2022-12-26	Payson, AZ	United States
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	United States
88888	20221230106513	Accident	ERA23LA097	2022-12-29	Athens, GA	United States


5 rows × 31 columns

In [15]:

```
# check the statistical summary of the data
df.describe()
```

Out[15]:

	Number.ofEngines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Tot
<b>count</b>	82805.000000	77488.000000	76379.000000	76956.000000	ε
<b>mean</b>	1.146585	0.647855	0.279881	0.357061	
<b>std</b>	0.446510	5.485960	1.544084	2.235625	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	0.000000	0.000000	0.000000	
<b>50%</b>	1.000000	0.000000	0.000000	0.000000	
<b>75%</b>	1.000000	0.000000	0.000000	0.000000	
<b>max</b>	8.000000	349.000000	161.000000	380.000000	



For the column to be relevant to the analysis, it should answer the questions:

- Which? Answer types of aircraft or models.
- Why? Answers what caused the incident
- When? Answer the time trend
- What kind of incident happened? Answer types of incidents and which one of them is sever

## 2. Data cleaning

- Handle missing values and duplicated values.
- Check on inconsistent formatting in the key fields.
- Standardize categorical variables for accurate grouping.
- Check on outliers.....

```
In [17]: # check on duplicated values
df.duplicated().any()
```

Out[17]: False

```
In [19]: # Check out the columns with missing values
df.isnull().any()
```

```
Out[19]: Event.Id                False
Investigation.Type              False
Accident.Number                 False
Event.Date                     False
Location                        True
Country                         True
Latitude                        True
Longitude                       True
Airport.Code                    True
Airport.Name                    True
Injury.Severity                 True
Aircraft.damage                 True
Aircraft.Category               True
Registration.Number             True
Make                           True
Model                           True
Amateur.Built                   True
Number.of.Engines               True
Engine.Type                     True
FAR.Description                 True
Schedule                        True
Purpose.of.flight               True
Air.carrier                     True
Total.Fatal.Injuries            True
Total.Serious.Injuries          True
Total.Minor.Injuries            True
Total.Uninjured                 True
Weather.Condition               True
Broad.phase.of.flight           True
Report.Status                   True
Publication.Date                True
dtype: bool
```

```
In [21]: # check the columns missing values with their percentages
df.isnull().mean().sort_values(ascending=False).round(2) * 100
```

```
Out[21]: Schedule      86.0
Air.carrier      81.0
FAR.Description   64.0
Aircraft.Category 64.0
Longitude        61.0
Latitude         61.0
Airport.Code     44.0
Airport.Name     41.0
Broad.phase.of.flight 31.0
Publication.Date 15.0
Total.Serious.Injuries 14.0
Total.Minor.Injuries 13.0
Total.Fatal.Injuries 13.0
Engine.Type      8.0
Report.Status    7.0
Purpose.of.flight 7.0
Number.of.Engines 7.0
Total.Uninjured  7.0
Weather.Condition 5.0
Aircraft.damage  4.0
Registration.Number 2.0
Injury.Severity  1.0
Country          0.0
Amateur.Built    0.0
Model            0.0
Make             0.0
Location         0.0
Investigation.Type 0.0
Event.Date       0.0
Accident.Number  0.0
Event.Id         0.0
dtype: float64
```

```
In [23]: # Drop the unnecessary columns, I mean according to the analysis
df.drop(['Schedule', 'Airport.Code', 'Registration.Number'], axis=1, inplace= True)
```

```
In [25]: # Replace columns with unknown values.
# This is to avoid biased data.
df['Air.carrier'] = df['Air.carrier'].fillna('Unknown Operator')
df['FAR.Description'] = df['FAR.Description'].fillna('Unknown FAR Category')
df['Aircraft.Category'] = df['Aircraft.Category'].fillna('Unknown Category')
df['Airport.Name'] = df['Airport.Name'].fillna('Unknown Airport')
df['Broad.phase.of.flight'] = df['Broad.phase.of.flight'].fillna('Unknown Phase')
```

```
In [27]: # check for the remaining columns with missing values .
df.isnull().sum().sort_values(ascending=False)
```

```
Out[27]: Longitude      54516
Latitude      54507
Publication.Date 13771
Total.Serious.Injuries 12510
Total.Minor.Injuries 11933
Total.Fatal.Injuries 11401
Engine.Type      7096
Report.Status    6384
Purpose.of.flight 6192
Number.of.Engines 6084
Total.Uninjured  5912
Weather.Condition 4492
Aircraft.damage  3194
Injury.Severity  1000
Country          226
Amateur.Built    102
Model            92
Make             63
Location         52
Investigation.Type 0
FAR.Description  0
Air.carrier      0
Aircraft.Category 0
Airport.Name     0
Event.Date       0
Broad.phase.of.flight 0
Accident.Number  0
Event.Id         0
dtype: int64
```

```
In [29]: # Lets check the mean and median of this columns
cols_int = [
    'Total.Serious.Injuries',
    'Total.Minor.Injuries',
    'Total.Fatal.Injuries',
    'Total.Uninjured',
    'Number.of.Engines'
]

# Check mean for all at once
df[cols_int].mean()

# check medians for the columns
df[cols_int].median()
```

```
Out[29]: Total.Serious.Injuries    0.0
Total.Minor.Injuries              0.0
Total.Fatal.Injuries              0.0
Total.Uninjured                  1.0
Number.of.Engines                 1.0
dtype: float64
```


So, according to the data above we cannot fill the missing values of the injured columns with the mean and median but we can fill them with 0



```
In [31]: # fill the missing values of the injured cols with 0
injury_cols = ['Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Fatal.Injuries']
df[injury_cols] = df[injury_cols].fillna(0)
```

```
In [33]: # Lets look into the number.of. Engines col using a visual to check if the data is
import matplotlib.pyplot as plt

df['Number.of.Engines'].dropna().plot(kind='hist', bins=10, color='green', ec='black')
plt.title("Distribution of Number of Engines")
plt.show()
```



```
In [35]: # fill the data with median
df['Number.of.Engines'] = df['Number.of.Engines'].fillna(df['Number.of.Engines'].median())
```

```
In [37]: # check for the remaining columns with missing values .
df.isnull().sum().sort_values(ascending=False)
```

```
Out[37]: Longitude      54516
          Latitude       54507
          Publication.Date 13771
          Engine.Type     7096
          Report.Status   6384
          Purpose.of.flight 6192
          Weather.Condition 4492
          Aircraft.damage  3194
          Injury.Severity  1000
          Country         226
          Amateur.Built   102
          Model           92
          Make            63
          Location        52
          Aircraft.Category 0
          Total.Serious.Injuries 0
          Accident.Number  0
          Broad.phase.of.flight 0
          Event.Date       0
          Total.Uninjured  0
          Total.Minor.Injuries 0
          Air.carrier      0
          Total.Fatal.Injuries 0
          FAR.Description  0
          Airport.Name     0
          Number.of.Engines 0
          Investigation.Type 0
          Event.Id         0
          dtype: int64
```

```
In [39]: # ckeck the injury.severity col
         df[df['Injury.Severity'].isnull()]
```

Out[39]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Coun
63918	20080111X00038	Incident	DCA08WA024	2008-01-03	Deauville Saint, France	Fra
63962	20080204X00132	Accident	NYC08WA081	2008-01-16	Kiteni, Peru	P
63987	20080304X00254	Incident	ENG08RA015	2008-01-24	Kingston, Jamaica	Jama
64026	20081219X65255	Incident	ENG08WA014	2008-02-03	Nurnberg, Germany	Germi
64128	20080409X00444	Accident	NYC08WA121	2008-02-28	Lago Ranco, Chile	Cl
...	...	...	...	...	...	...
88863	20221213106449	Accident	GAA22WA311	2022-12-11	Kildare,	Irela
88874	20221215106462	Accident	CEN23LA064	2022-12-15	Patterson, LA	Uni Sta
88879	20221219106472	Accident	DCA23LA096	2022-12-18	Kahului, HI	Uni Sta
88885	20221227106494	Accident	ERA23LA095	2022-12-26	Hampton, NH	Uni Sta
88887	20221227106498	Accident	WPR23LA076	2022-12-26	Morgan, UT	Uni Sta

1000 rows × 28 columns



```
In [41]: # We can fill up this column with the injury columns .

def refer_severity(row):
    if row['Total.Fatal.Injuries'] > 0:
        return 'Fatal'
    elif row['Total.Serious.Injuries'] > 0:
        return 'Serious'
    elif row['Total.Minor.Injuries'] > 0:
        return 'Minor'
    elif row['Total.Uninjured'] > 0:
        return 'None'
    else:
        return 'Unknown'

df['Injury.Severity'] = df.apply(refer_severity, axis=1)
```

```
In [43]: df['Injury.Severity'].value_counts()
```

```
Out[43]: Injury.Severity
None      47089
Fatal      17813
Minor      11488
Serious    11190
Unknown     1309
Name: count, dtype: int64
```

```
In [45]: # check the country and location
df[df['Country'].isnull()]
```

Out[45]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location
36	20020917X02410	Accident	MIA82FKA05	1982-01-04	SAINT CROIX
464	20020917X02358	Accident	MIA82DA062	1982-03-02	HUMA CAO
465	20020917X02026	Accident	FTW82DA076	1982-03-02	MUSTANG BLK A11
725	20020917X02377	Accident	MIA82DA091	1982-03-31	MOCA
831	20020917X02069	Accident	FTW82DA127	1982-04-13	WEST DELTA 105D
...	...	...	...	...	...
52288	20020322X00387	Accident	DCA02MA029	2002-03-22	int'l waters
54284	20040528X00699	Accident	DCA03WA031	2003-03-12	JOHANNESBURG
56200	20040528X00697	Accident	DCA04WA026	2004-02-10	Sharjah Airport
58803	20050616X00790	Incident	DCA05WA073	2005-06-10	Canada/US borde
62530	20070518X00582	Incident	DCA07WA043	2007-05-13	London Control

226 rows × 28 columns



In [47]: `# check on the remaining cols with missing values.  
df.isnull().sum().sort_values(ascending=False)`

```
Out[47]: Longitude      54516
          Latitude      54507
          Publication.Date 13771
          Engine.Type    7096
          Report.Status  6384
          Purpose.of.flight 6192
          Weather.Condition 4492
          Aircraft.damage 3194
          Country        226
          Amateur.Built  102
          Model          92
          Make           63
          Location       52
          Injury.Severity 0
          Total.Serious.Injuries 0
          Accident.Number 0
          Broad.phase.of.flight 0
          Event.Date     0
          Total.Uninjured 0
          Total.Minor.Injuries 0
          Air.carrier     0
          Total.Fatal.Injuries 0
          Aircraft.Category 0
          FAR.Description 0
          Airport.Name    0
          Number.of.Engines 0
          Investigation.Type 0
          Event.Id        0
          dtype: int64
```

```
In [49]: #Fill the remaining columns with unknowns.
          # To prevent data irrelevance or bias.
          rem_cols=['Engine.Type', 'Report.Status', 'Purpose.of.flight', 'Weather.Condition', 'Ai
                  'Country', 'Amateur.Built', 'Model', 'Make', 'Location']
          df[rem_cols]= df[rem_cols].fillna('Unknown')
```

```
In [51]: # check on the remaining cols with missing values.
          df.isnull().sum().sort_values(ascending=False)
```

```
Out[51]: Longitude      54516
Latitude      54507
Publication.Date 13771
Number.of.Engines      0
Report.Status      0
Broad.phase.of.flight  0
Weather.Condition      0
Total.Uninjured      0
Total.Minor.Injuries   0
Total.Serious.Injuries 0
Total.Fatal.Injuries   0
Air.carrier      0
Purpose.of.flight      0
FAR.Description      0
Engine.Type      0
Event.Id      0
Investigation.Type      0
Model      0
Make      0
Aircraft.Category      0
Aircraft.damage      0
Injury.Severity      0
Airport.Name      0
Country      0
Location      0
Event.Date      0
Accident.Number      0
Amateur.Built      0
dtype: int64
```

```
In [53]: # Lets clean the latitude and longitude columns for Tableau mappings
#Ensure 'Latitude' and 'Longitude' are numeric
df['Latitude'] = pd.to_numeric(df['Latitude'], errors='coerce')
df['Longitude'] = pd.to_numeric(df['Longitude'], errors='coerce')
```

```
In [55]: # Drop rows without coordinates needed for mapping
df = df.dropna(subset=['Latitude', 'Longitude'])
```

```
In [57]: # Start fresh with a clean copy if you filtered before
df = df.copy()

#apply your parsing safely
df['Publication.Date'] = pd.to_datetime(df['Publication.Date'], format='%d-%m-%Y',
```

```
In [59]: #create a display version that fills missing with 'Unknown'
df['PubDate_Display'] = df['Publication.Date'].dt.strftime('%Y-%m-%d')
df['PubDate_Display'] = df['PubDate_Display'].fillna('Unknown')
```

```
In [61]: # counter check if the data is ready for use.
df.head(15)
```

Out[61]:

	Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Cou
2	20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	U S
5	20170710X52551	Accident	NYC79AA106	1979-09-17	BOSTON, MA	U S
593	20080417X00504	Accident	MIA08CA076	1982-03-16	MOBILE, AL	U S
3654	20051208X01953	Accident	SEA83LA209	1983-01-08	Goldendale, WA	U S
6202	20020904X01525	Accident	SEA83FA208	1983-09-09	Kalispell, MT	U S
22096	20001213X27446	Accident	LAX89LA068	1988-12-23	Midway Islands, PO	U S
24567	20021022X05356	Accident	CHI90LA280	1989-12-01	ENGADINE, MI	U S
26826	20030411X00484	Accident	ANC91GAMS1	1990-10-11	Deadhorse, AK	U S
31353	20170710X10920	Accident	FTW92FA224	1992-09-05	Alpine, TX	U S
38740	20011127X02295	Accident	NYC96FA192	1995-11-28	Marlinton, WV	U S
42691	20001208X08803	Accident	CHI97FA308	1997-09-14	St. Ignaces, MI	U S
44870	20001211X11043	Accident	FTW98FA380	1998-09-11	HOUSTON, TX	U S
45203	20041203X01907	Accident	ATL99FA136	1998-11-04	Robbinsville, NC	U S
45404	20001211X11573	Accident	LAX99FA051	1998-12-17	LOS ANGELES, CA	U S
45592	20001205X00119	Incident	ANC99IA027	1999-02-05	FAIRBANKS, AK	U S

15 rows × 29 columns



```
In [63]: # check on outliers on the numeric columns
import seaborn as sns
# Select numeric columns
```



```

numeric_cols = df.select_dtypes(include='number').columns

# Set the number of rows and columns for the subplot grid
nrows = 2
ncols = 3

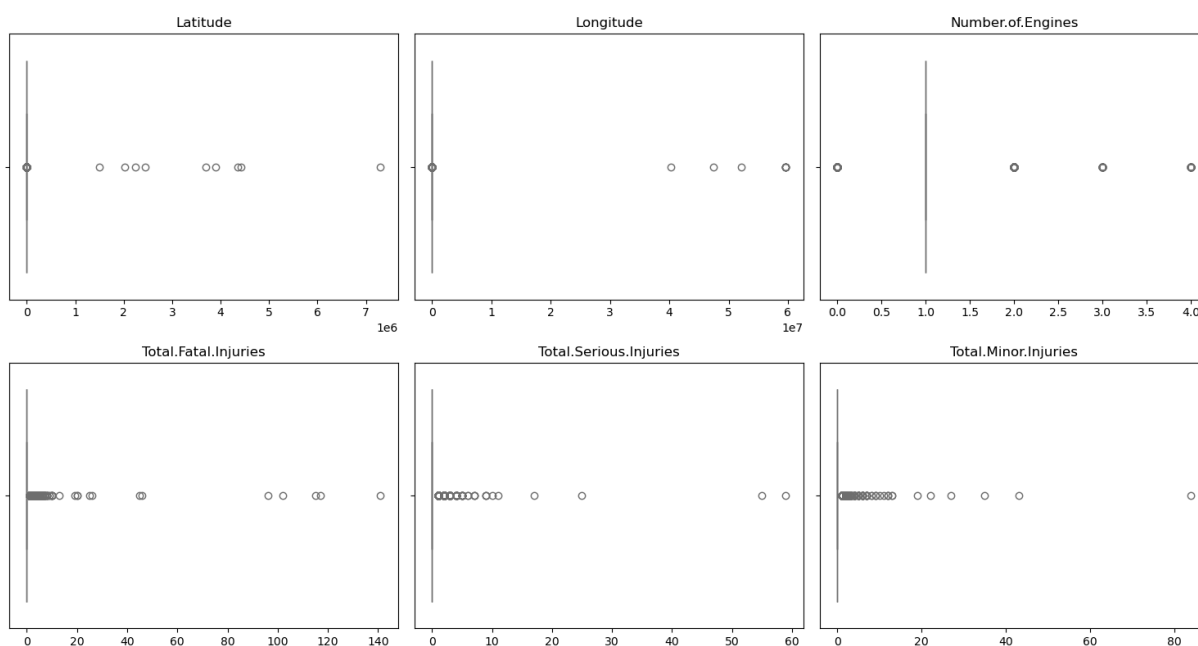
# Create subplots
fig, axes = plt.subplots(nrows=nrows, ncols=ncols, figsize=(15, 8))
axes = axes.flatten() # Flatten for easy indexing

# Plot each numeric column
for i, col in enumerate(numeric_cols[:nrows * ncols]): # Limit to fit grid
    sns.boxplot(data=df, x=col, ax=axes[i], color='skyblue')
    axes[i].set_title(f'{col}')
    axes[i].set_xlabel('')

# Hide any unused subplots
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



### 3. Univariate & Bivariate Analysis

- Explore the most common aircraft types and count how many incidents each has.
- Analyze incident trends over time to detect rising or declining risk patterns

```

In [64]: # check the aircraft type and count how many incidents each has
df['Model'].value_counts().head()

```

```
Out[64]: Model
          152      195
          172N     175
          172S     154
          172      122
          PA-28-140 117
          Name: count, dtype: int64
```

```
In [67]: # Analyze incident trends over time to detect rising or declining risk patterns
# Group the model data with the event year column.
# Change the event date column to real-time dates
df['Event.Date']=pd.to_datetime(df['Event.Date'], errors= 'coerce')

# extract the years from the `df['Event.Date']`
df['Event.Year']= df['Event.Date'].dt.year
```

```
In [69]: grouped_df=df.groupby(['Event.Year', 'Model']).size().reset_index(name= 'incident_tt')
```


```
In [71]: top_models = df['Model'].value_counts().head().index
grouped_df = grouped_df[grouped_df['Model'].isin(top_models)]
grouped_df
```

Out[71]:

	Event.Year	Model	incident_ttl
<b>14</b>	1999	152	2
<b>27</b>	2000	172	1
<b>30</b>	2000	172N	2
<b>32</b>	2000	172S	1
<b>127</b>	2000	PA-28-140	3
<b>181</b>	2001	152	28
<b>185</b>	2001	172	11
<b>193</b>	2001	172N	11
<b>197</b>	2001	172S	8
<b>548</b>	2001	PA-28-140	4
<b>703</b>	2002	152	31
<b>713</b>	2002	172	19
<b>725</b>	2002	172N	35
<b>730</b>	2002	172S	18
<b>1287</b>	2002	PA-28-140	20
<b>1548</b>	2003	152	29
<b>1555</b>	2003	172	14
<b>1568</b>	2003	172N	29
<b>1572</b>	2003	172S	29
<b>2214</b>	2003	PA-28-140	18
<b>2519</b>	2004	152	19
<b>2528</b>	2004	172	17
<b>2541</b>	2004	172N	33
<b>2545</b>	2004	172S	23
<b>3133</b>	2004	PA-28-140	14
<b>3396</b>	2005	152	36
<b>3405</b>	2005	172	27
<b>3418</b>	2005	172N	23
<b>3422</b>	2005	172S	26
<b>4039</b>	2005	PA-28-140	19

	Event.Year	Model	incident_ttl
4319	2006	152	26
4326	2006	172	13
4341	2006	172N	22
4345	2006	172S	22
4948	2006	PA-28-140	19
5208	2007	152	24
5215	2007	172	20
5231	2007	172N	20
5235	2007	172S	27
5872	2007	PA-28-140	20

```
In [73]: # plot the data for better visualization
plt.figure(figsize=(12, 6))
sns.lineplot(data=grouped_df, x='Event.Year', y='incident_ttl', hue='Model', marker)
plt.title('Incident Trends Over Time by Aircraft Model')
plt.xlabel('Year')
plt.ylabel('Number of Incidents')
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
In [75]: grouped = df.groupby(['Event.Year', 'Model', 'Injury.Severity']).size().reset_index
grouped
```

Out[75]:

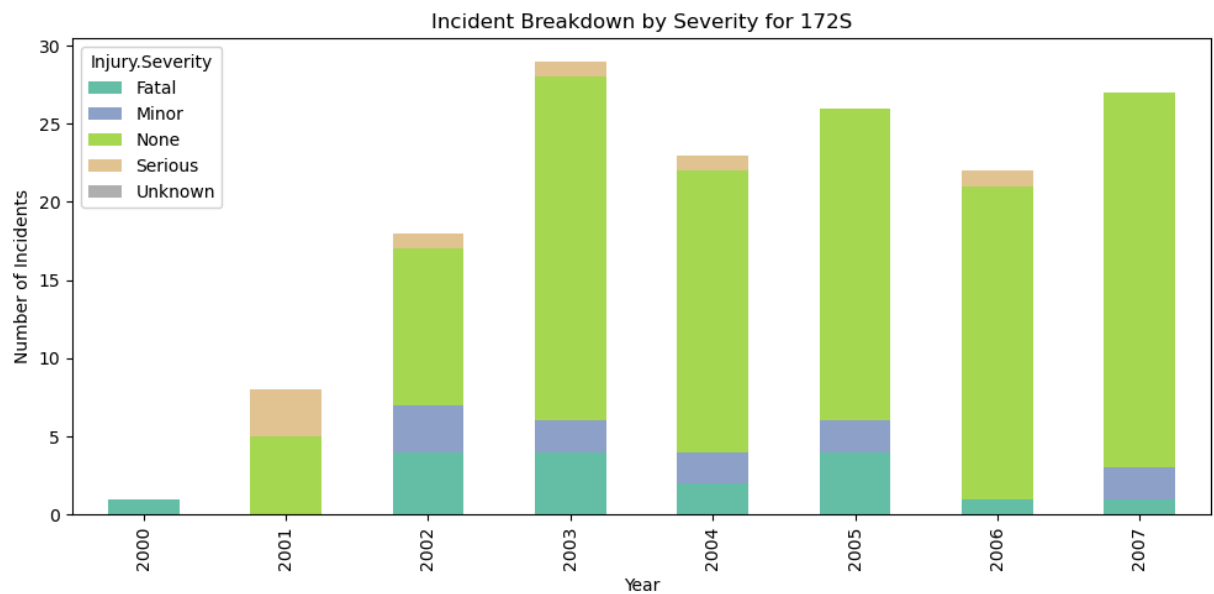


	Event.Year	Model	Injury.Severity	Incident_ttl
0	1974	172M	Fatal	1
1	1979	DC9	Minor	1
2	1982	C24R	Fatal	1
3	1983	14-19-3	Fatal	1
4	1983	182N	Fatal	1
...	...	...	...	...
8149	2021	T303	Fatal	1
8150	2021	U206	None	1
8151	2022	182	None	1
8152	2022	M-7-235B	None	1
8153	2022	PA-22-160	None	1

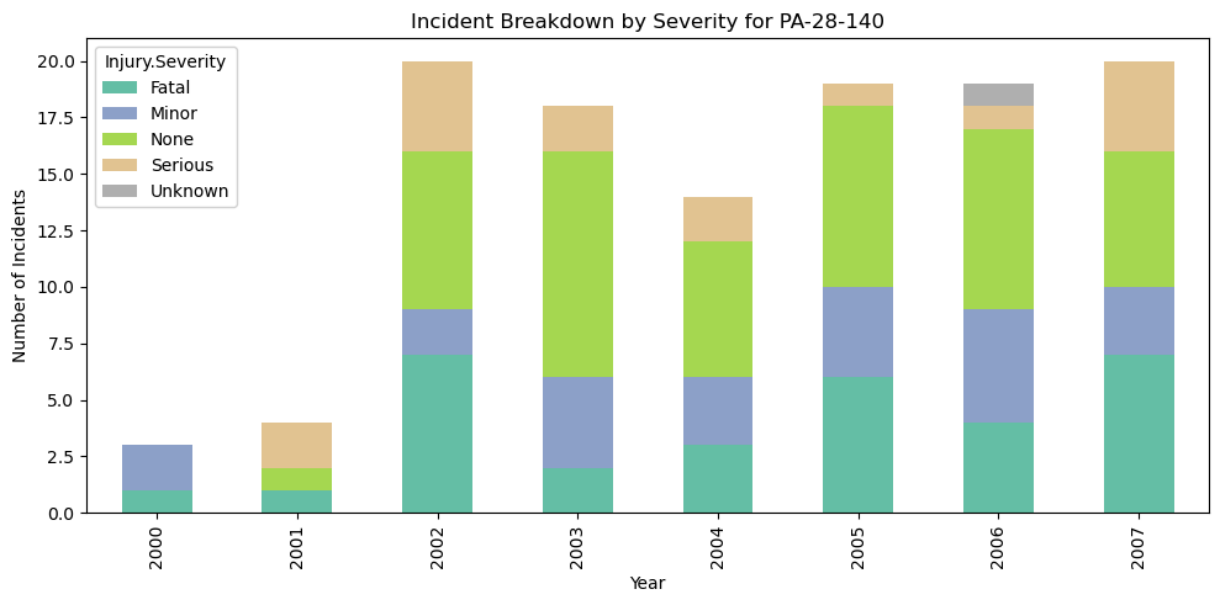

8154 rows × 4 columns

```
In [77]: # Change this grouped data to a pivot table to be able to plot a grouped bar format
pivot_df = grouped.pivot_table(index=['Event.Year', 'Model'],
                                columns='Injury.Severity',
                                values='Incident_ttl',
                                fill_value=0).reset_index()
```

```
In [79]: # Loop through the top 5 models and plot for each
for model_name in top_models:
    model_data = pivot_df[pivot_df['Model'] == model_name]

    # Plot a stacked bar chart
    model_data.plot(x='Event.Year', kind='bar', stacked=True, figsize=(10, 5), color=
plt.title(f'Incident Breakdown by Severity for {model_name}')
plt.ylabel('Number of Incidents')
plt.xlabel('Year')
plt.tight_layout()
plt.show()
```





findings.

- The PA-28-140 Aircraft model has the highest constant incidents. It shows a high probability of danger, while the 172S model shows a minimal probability of danger.
- For further model analysis, will conduct on tableau for more insights.

#### 4. Multivariate & Grouped Insights

Compare incident frequencies by usage type (commercial vs private), age, or manufacturer.

*Changing the aircraft type from `Model` to `Aircraft` category for more simple and easy visualization creations.*

```
In [81]: # retrieve the columns to be used
df['Aircraft.Category'].value_counts()
```

```
Out[81]: Aircraft.Category
Unknown Category      8334
Airplane              3084
Helicopter            390
Glider                88
Balloon               40
Gyrocraft             21
Ultralight            6
Blimp                 3
Powered-Lift          2
Name: count, dtype: int64
```

```
In [83]: # can be used for advanced analysis of this objective .
df['Purpose.of.flight'].value_counts()
```

```
Out[83]: Purpose.of.flight
Personal              7044
Instructional         1559
Unknown              1188
Aerial Application    483
Business              403
Positioning           374
Other Work Use        241
Public Aircraft       149
Flight Test           124
Aerial Observation     94
Executive/corporate    59
Ferry                 58
Air Race/show         57
Skydiving             31
Public Aircraft - Federal 26
Banner Tow            23
External Load         17
Public Aircraft - State 14
Glider Tow            10
Public Aircraft - Local 9
Firefighting          4
Air Drop              1
Name: count, dtype: int64
```

```
In [85]: df['FAR.Description'].value_counts()
```



```
Out[85]: FAR.Description
Unknown FAR Category      8330
Part 91: General Aviation  3244
Part 137: Agricultural    162
Part 135: Air Taxi & Commuter  85
Part 121: Air Carrier     44
Non-U.S., Non-Commercial  35
Non-U.S., Commercial      22
Public Use                18
Unknown                   6
Part 129: Foreign         5
Part 133: Rotorcraft Ext. Load  5
NUSN                      5
091                       3
Part 91 Subpart K: Fractional  1
Part 125: 20+ Pax,6000+ lbs  1
Public Aircraft           1
NUSC                      1
Name: count, dtype: int64
```

```
In [87]: df['Make'].value_counts()
```

```
Out[87]: Make
Cessna      3485
Piper       1854
Beech       751
Bell        361
Robinson    308
...
Davenport (van's)  1
Dantzer Lawrence L  1
Aircraft Mfg & Development Co.  1
Murphy          1
MAULE           1
Name: count, Length: 1715, dtype: int64
```

```
In [89]: # Calculate the Aircraft age.
# Limitation: there is no manufacturing Date or Certificate.Issued.Date
# For temporary trends, use the `Event.Year` to analyze usage .
```

```
In [91]: # use the df['FAR.Description'] to create a usage_type column
def classify_usage(far):
    if pd.isna(far):
        return 'Unknown'
    elif '121' in far or '135' in far:
        return 'Commercial'
    elif '91' in far:
        return 'Private'
    else:
        return 'Other'

df['Usage.Type'] = df['FAR.Description'].apply(classify_usage)
```

```
In [93]: # create a pivot table to create incidents by Aircraft Category and usage.
pivot_t1=df.pivot_table(index= 'Aircraft.Category',
```

```

        columns= 'Usage.Type',
        values= 'Event.Id',
        aggfunc='count',
        fill_value=0
    )
pivot_t1

```

Out[93]:

	Usage.Type	Commercial	Other	Private
Aircraft.Category				
Airplane		102	218	2764
Balloon		0	0	40
Blimp		0	0	3
Glider		0	1	87
Gyrocraft		0	0	21
Helicopter		27	42	321
Powered-Lift		0	0	2
Ultralight		0	0	6
Unknown Category		0	8330	4

```

In [95]: # create a pivot table to gain incidents by Manufacturer and usage_type
pivot_t2=df.pivot_table(index='Make',
        columns='Usage.Type',
        values='Event.Id',
        aggfunc='count',
        fill_value=0
    ).sort_values(by='Commercial', ascending=False) # arranges by descending order us
pivot_t2

```

Out[95]:

Usage.Type	Commercial	Other	Private
Make			
Cessna	28	2365	1092
Boeing	16	187	17
Beech	9	597	145
Bell	9	288	64
Mcdonnell Douglas	9	72	3
...	...	...	...
Fox Alfred C.	0	1	0
Found Aircraft Canada Inc	0	0	2
Found Aircraft Canada	0	1	0
Found Acft	0	0	1
Zorn	0	1	0

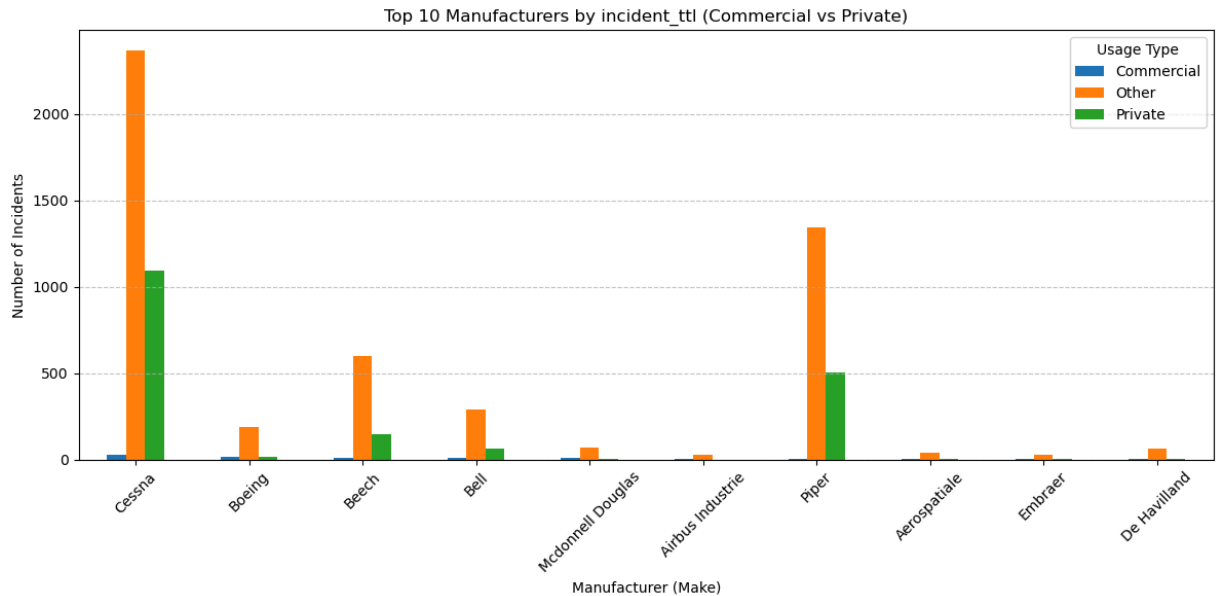
1715 rows × 3 columns

```

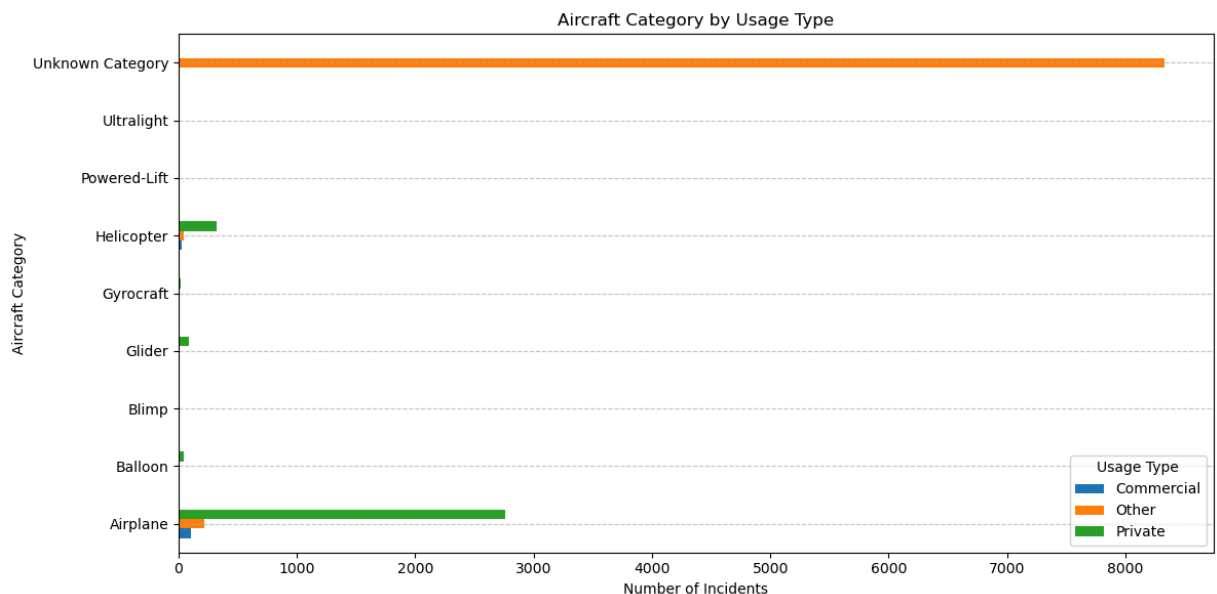
In [97]: # plot the pivot_t2 to check on the top ten manufacturers for the plane models.
Top_10=pivot_t2.head(10)
# plot a grouped bar chart
Top_10.plot(kind='bar', figsize=(12, 6))

# Label the chart
plt.title('Top 10 Manufacturers by incident_ttl (Commercial vs Private)')
plt.ylabel('Number of Incidents')
plt.xlabel('Manufacturer (Make)')
plt.xticks(rotation=45)
plt.legend(title='Usage Type')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```

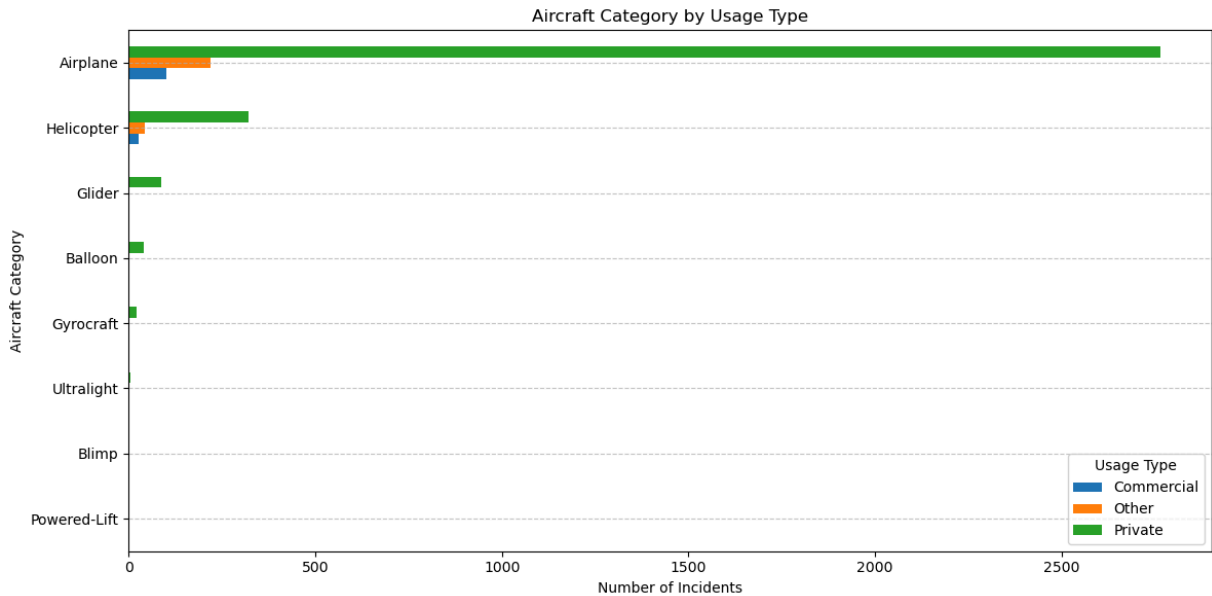


```
In [99]: # plot the pivot_t1 to check on the top Aircraft category by usage type.
# plot a grouped barh graph.
pivot_t1.plot(kind='barh', figsize=(12, 6))
# Label chart
plt.title('Aircraft Category by Usage Type')
plt.xlabel('Number of Incidents')
plt.ylabel('Aircraft Category')
plt.legend(title='Usage Type')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



```
In [101... # improved graph
# remove the Unknown category and sort the values to ascending order
pivot_filtered=pivot_t1.drop(index= 'Unknown Category', errors= 'ignore')
pivot = pivot_filtered.sort_values(by=pivot_filtered.columns.to_list(), ascending=True)
# plot the data.
pivot.plot(kind='barh', figsize=(12, 6))
```

```
# Label chart
plt.title('Aircraft Category by Usage Type')
plt.xlabel('Number of Incidents')
plt.ylabel('Aircraft Category')
plt.legend(title='Usage Type', loc=4)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



### My Findings

- according to the visualizations, private aviation has more incidents than the commercial ones.
- The manufacturers(Make) cessna and piper have high incidents do to their products especially on the private aircrafts

### 5. Incident Causes & Locations

- Visualize the top causes of incidents and the regions with the highest incident density.

In [103... `df.columns`

Out[103... `Index(['Event.Id', 'Investigation.Type', 'Accident.Number', 'Event.Date', 'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Name', 'Injury.Severity', 'Aircraft.damage', 'Aircraft.Category', 'Make', 'Model', 'Amateur.Built', 'Number.ofEngines', 'Engine.Type', 'FAR.Description', 'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status', 'Publication.Date', 'PubDate_Display', 'Event.Year', 'Usage.Type'], dtype='object')`

In [105... `# Let's load the recommended columns`  
`df['Broad.phase.of.flight'].value_counts()`

```
Out[105... Broad.phase.of.flight
Landing      3215
Takeoff      2195
Cruise       1775
Maneuvering   1475
Approach      1221
Unknown Phase  483
Taxi          377
Descent       343
Climb         339
Go-around     263
Standing      199
Unknown        65
Other         18
Name: count, dtype: int64
```

```
In [107... df['Injury.Severity'].value_counts()
```

```
Out[107... Injury.Severity
None      6337
Fatal     2621
Minor     1716
Serious   1287
Unknown      7
Name: count, dtype: int64
```

```
In [109... # Group and reshape your clean data
phase_severity = df.groupby(['Broad.phase.of.flight', 'Injury.Severity']).size().un
phase_severity
```

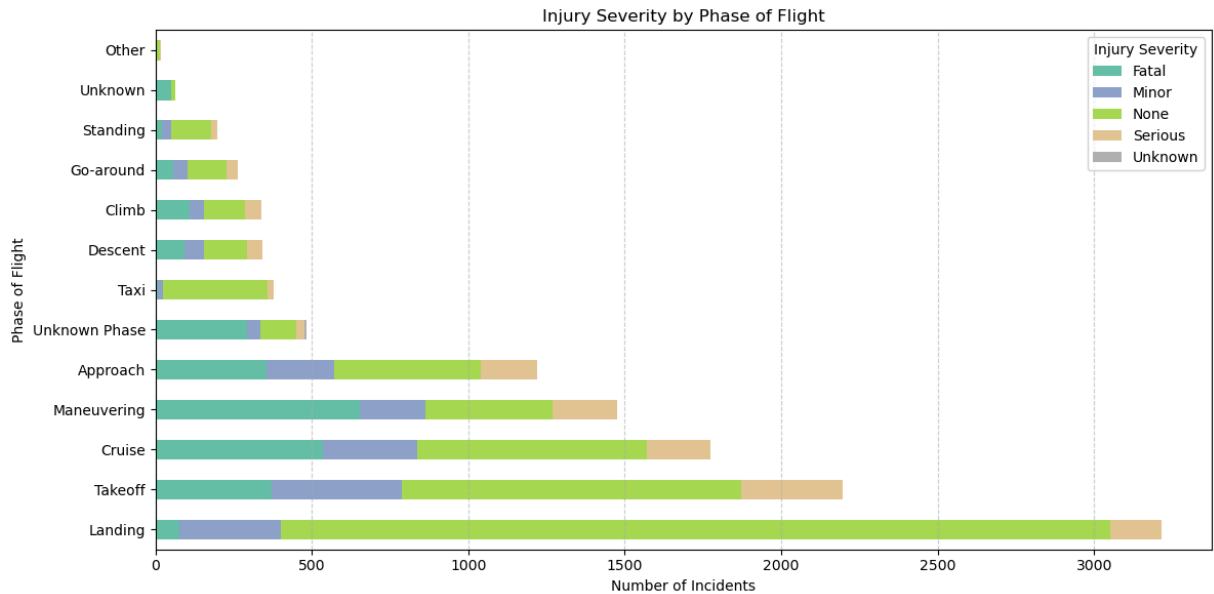
Out[109...

Injury.Severity	Fatal	Minor	None	Serious	Unknown
Broad.phase.of.flight					
Approach	356	214	471	180	0
Climb	108	48	129	54	0
Cruise	534	302	735	204	0
Descent	92	63	139	49	0
Go-around	56	46	126	35	0
Landing	78	322	2650	165	0
Maneuvering	653	210	405	207	0
Other	3	3	8	4	0
Standing	21	28	128	22	0
Takeoff	370	419	1084	322	0
Taxi	7	18	333	19	0
Unknown	51	0	14	0	0
Unknown Phase	292	43	115	26	7

In [111...

```
# Let's plot the two columns to plot a grouped bar
# Sort by total incident count descending
phase_severity = phase_severity.loc[phase_severity.sum(axis=1).sort_values(ascending=False)]
# Plot
phase_severity.plot(kind='barh', stacked=True, figsize=(12, 6), colormap='Set2')

# Labels and formatting
plt.title('Injury Severity by Phase of Flight')
plt.xlabel('Number of Incidents')
plt.ylabel('Phase of Flight')
plt.legend(title='Injury Severity')
plt.grid(axis='x', linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



### My Finding

- The phase of flight with high incidents is landing.

## 5. Insights and Recommendations

- Other and Private usage types of aircraft have higher injury severity than the commercial ones. This suggests that general aviation poses a higher volume of safety incidents, possibly due to looser regulations, pilot experience, or maintenance variation.
- The Aircraft Category poses high incidents, most especially in private/personal aviation. However, the multi-engine aircraft categories, such as Gyrocraft, Powered-Lift, and Ultralight, have had slight incidents, but these aircraft can be suitable for starting a business venture in aviation.
- While the landing phase accounts for the highest number of incidents, most result in non-injury outcomes. In contrast, the maneuvering, takeoff, and cruise phases—though less frequent—are associated with high injury severity. This suggests a need to shift to a safety focus on in-flight operations, where risk per incident may be higher.

## 6. Challenges and conclusions

- The data had a lot of critical missing values.

```
In [115... # Lets change the data back to csv for Tableau analysis.
df.to_csv('Aviation_clean_data.csv', index=False)
```